
Distractor Filtering Through Open Knowledge Graph Concept Proximity for Dense RAG Systems

June 9, 2026

Gianluca Como (1605533)

Abstract

Dense retrieval can surface *distractors* — passages close to the query but answer-free — that degrade RAG accuracy, as Cuconasu et al. [3] show. We test whether *pure graph topology* on Wikidata can filter them, without using relation semantics. We build a 661.5M-edge entity graph from a local HDT dump, score query–document pairs by topological reachability (connected $\times$ purity), and fuse this with the dense score across a grid of hop-distance, hub-threshold and mixing weight α , evaluated on 1000 NQ-open queries with an LLM judge. The result is a clean *negative*: the KG signal never beats the dense baseline (0.364) beyond noise (best 0.372) and actively degrades accuracy as its weight grows (down to 0.290); 13/90 conditions are significant under McNemar+Bonferroni, all of them worse. We explain the mechanism (reachability saturates on a dense graph) and document the engineering and the methodological departures from [3].

1. Introduction

Retrieval-Augmented Generation (RAG) prepends retrieved passages to an LLM’s prompt, but [3] show that high-scoring passages are often *distractors* — semantically similar to the query yet answer-free — and hurt the reader. We ask whether the topology of an open knowledge graph (Wikidata) can separate relevant passages from distractors. The hypothesis: a passage whose entities are *graph-close* to the query entities is less likely to be a distractor than one chosen by semantic similarity alone. Unlike approaches that exploit the graph’s full relation semantics, we test the weakest signal — bare proximity within k hops — to isolate its contribution.

Our contributions are: (i) a pipeline that navigates the *en-*

tire Wikidata graph from a local HDT dump and tests ≤ 3 -hop reachability meet-in-the-middle, without ever expanding a high-degree hub’s neighbour list; (ii) a clean negative result, supported by exact McNemar tests and a mechanistic explanation; (iii) three deliberate deviations from the evaluation setup of [3] (passage padding, an extraction prompt for the base model, an LLM judge), documented for reproducibility. Code: {4}.

2. Related work

We adopt the corpus, retriever and reader of [3], who characterise the distractor problem in RAG. Retrieval uses Contriever over the gold-augmented Wikipedia corpus of [3]; entity linking uses the ReFinED linker of Ayoola et al. [1], one of the systems evaluated by Bast et al. [2]. Knowledge-graph approaches to RAG use the graph’s relations and semantics; we instead use only its topology. The graph is stored as HDT over a Wikidata dump.

3. Method

Data and retrieval. The corpus is the gold-augmented Dec-2018 Wikipedia release of [3]{1}, Contriever-encoded and indexed with FAISS (IndexFlatIP, ~ 42 M passages). From NQ-open {2} we keep queries with short (≤ 5 -token) answers — matching the extractive setup of [3] — where the question and *all* answer variants link to Wikidata entities (ReFinED), and evaluate on a curated 1000-query subset.

Passage segmentation (departure 1). The corpus of [3] ships as raw 100-word passages with ragged, unpadding tails. We instead cut on sentence boundaries and pad each segment back to exactly 100 words using the article’s opening words. This avoids subject-less chunks (e.g. a passage starting mid-sentence) that break entity linking and reader comprehension, keeps passage length uniform, and incidentally anchors every passage to the article subject. Contriever mean-pools, so fixed length is a consistency choice, not a model requirement.

Email: Gianluca Como <comogianluca@gmail.com>.

Graph construction (engineering). We query Wikidata from a pre-built HDT dump {3} (Jan 2022, ~166 GB) via pyHDT; the public SPARQL endpoint was abandoned after even a minimal COUNT on a hub (Q5, *human*) timed out. We export every direct entity-to-entity statement (Wikidata’s “truthy” wdt: predicates), one file per predicate, yielding 6.61×10^8 edges. For a pair (q, d) we decide $\text{dist}(q, d) \leq 3$ meet-in-the-middle: distance 1 is a set lookup, distance 2 an intersection of 1-hop neighbourhoods N_1 , distance 3 an indexed edge-probe between $N_1(q)$ and $N_1(d)$. We never materialise the neighbourhood of a hub (Q30, *USA*, has 2.35M edges); we only test membership. A degree threshold t optionally drops high-degree bridge nodes.

Scoring and fusion. Following the proposal, the KG score is $s_{\text{kg}} = \text{cr} \cdot \text{pr}$, where the *connected ratio* cr is the fraction of query entities reaching ≥ 1 document entity within k hops and the *purity ratio* pr the fraction of document entities reached by any query entity. We fuse with the dense score as

$$s = (1 - \alpha) s_{\text{dense}} + \alpha s_{\text{kg}}, \quad (1)$$

with α the *KG weight* and $s_{\text{dense}} = \sigma / \max_q \sigma$ a per-query max-only rescaling of the raw Contriever score σ that caps it at 1 without stretching its narrow band (min-max would). We sweep a grid of distance $k \in \{1, 2, 3\}$, threshold $t \in \{500, 1k, 2k, 5k, 10k, \infty\}$ and $\alpha \in \{0.1, 0.3, 0.5, 0.7, 0.9\}$: 90 rerank cells plus the dense baseline.

Generation and evaluation (departures 2–3). The top-5 passages feed Llama-2-7B (base). The prompt is derived from [3] with two deviations: a one-shot exemplar (the base, non-instruct model otherwise continues “Answer:” with Wikipedia-style prose) and the question placed *above* the documents (the paper places it after). Correctness is decided by an *LLM judge* (Qwen2.5-7B-Instruct, a different model to avoid self-preference) that checks whether the response contains one of the NQ-open gold answers. This replaces the brittle substring exact-match of [3], which misses paraphrases and formatting.

4. Results

Baseline is sound. Dense retrieval scores 0.364, in line with the ≈ 0.36 that [3] report for Llama-2 with Contriever-retrieved documents on NQ-open — the harness is healthy and the comparison valid.

No gain, then harm. The best cell ($\alpha=0.3$, dist 2, $t=5k$) reaches 0.372: $+0.8\text{pp} \approx 8$ queries, i.e. noise (Table 1). As α grows the KG signal dominates and accuracy falls monotonically, bottoming at 0.290 ($\alpha=0.9$, dist 3, $t=5k$).

Table 1. LLM-judge accuracy, all 90 rerank settings (1000 NQ-open queries). Dense baseline = 0.364. Columns: KG weight α ; rows: hop-distance k , hub-threshold t (∞ =no filter; distance 1 has no bridge, so it is threshold-invariant). Best **bold**, worst underlined.

k	t	$\alpha=.1$	.3	.5	.7	.9
1	any	.353	.344	.343	.316	.322
2	500	.362	.358	.361	.348	.354
2	1,000	.364	.361	.358	.343	.350
2	2,000	.360	.371	.357	.337	.341
2	5,000	.357	.372	.362	.337	.337
2	10,000	.364	.365	.357	.334	.329
2	∞	.354	.343	.311	.298	.295
3	500	.354	.344	.341	.329	.335
3	1,000	.346	.355	.342	.322	.310
3	2,000	.354	.349	.316	.303	.293
3	5,000	.366	.351	.310	.292	<u>.290</u>
3	10,000	.361	.362	.321	.306	.304
3	∞	.347	.333	.314	.314	.318

Threshold gradient is the mechanism. Holding $\alpha=0.9$, dist 2, accuracy decays monotonically as the hub filter loosens, from 0.354 ($t=500$) to 0.295 ($t=\infty$): admitting high-degree nodes as bridges connects everything to everything and destroys the signal.

Significance. Exact McNemar tests against the baseline, Bonferroni-corrected for the 90 comparisons (significance threshold $0.05/90 \approx 5.6 \times 10^{-4}$), flag 13 of the 90 conditions as significant — *all* worse, none better; each flips $\approx 200/1000$ answers yet nets negative.

5. Discussion and conclusions

Why it fails. On a graph as dense as Wikidata, 2–3-hop reachability *saturates*: almost every document entity is reachable from almost every query entity, so $\text{cr} \approx 1$ and the score stops discriminating; the hub filter limits the damage but recovers no signal. For about 17% of queries the question entity is itself a very high-degree node (a country, a religion), where any passage is graph-close and the KG score is uninformative.

Limitations and future work. Our metric collapses relation semantics (P31 *instance-of* and P19 *place-of-birth* count alike), so it cannot screen out semantically irrelevant links. The honest next step is to use the graph’s relations and semantics — not just its bare topology.

Conclusion. Pure topological proximity does not filter distractors for dense RAG on NQ-open; the value is the explained negative result and the reusable Wikidata tooling.

Data and code

{1} Corpus: https://huggingface.co/datasets/florin-hf/wiki_dump2018_nq_open.

{2} Queries: https://huggingface.co/datasets/florin-hf/nq_open_gold.
 {3} Wikidata HDT dump: <https://hdt-dumps.cluster.ai.wu.ac.at/dumps/>.
 {4} Code: delivery repository (to be created).

References

- [1] Ayoola, T., Tyagi, S., Fisher, J., Christodoulopoulos, C., and Pierleoni, A. ReFinED: An efficient zero-shot-capable approach to end-to-end entity linking. In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics (NAACL): Industry Track*, 2022.
- [2] Bast, H., Hertel, M., and Prange, N. A fair and in-depth evaluation of existing end-to-end entity linking systems. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2023.
- [3] Cuconasu, F., Trappolini, G., Siciliano, F., Filice, S., Campagnano, C., Maarek, Y., Tonellotto, N., and Silvestri, F. The power of noise: Redefining retrieval for RAG systems. In *Proceedings of the 47th International ACM SIGIR Conference on Research and Development in Information Retrieval (SIGIR)*, 2024.
- [4] Naveed, H., Khan, A. U., Qiu, S., Saqib, M., Anwar, S., Usman, M., Akhtar, N., Barnes, N., and Mian, A. A comprehensive overview of large language models. *arXiv preprint arXiv:2307.06435*, 2023.
- [5] Peng, B., Zhu, Y., Liu, Y., Bo, X., Shi, H., Hong, C., Zhang, Y., and Tang, S. Graph retrieval-augmented generation: A survey. *arXiv preprint arXiv:2408.08921*, 2024.