
Distractor Filtering Through Open Knowledge Graph Concept Proximity for Dense RAG Systems

June 9, 2026

Gianluca Como (1605533)

Abstract

Dense retrieval can surface *distractors* — passages close to the query but answer-free — that degrade RAG accuracy, as Cuconasu et al. [3] show. We test whether *pure graph topology* on Wikidata can filter them, without using relation semantics. We build a 661.5M-edge entity graph from a local HDT dump, score query–document pairs by topological reachability (connected $\times$ purity), and fuse this with the dense score across a grid of hop-distance, hub-threshold and mixing weight α , evaluated on 1000 NQ-open queries with an LLM judge. The result is a clean *negative*: the KG signal never beats the dense baseline (0.364) beyond noise (best 0.372) and actively degrades accuracy as its weight grows (down to 0.290); 13/90 conditions are significant under McNemar+Bonferroni, all of them worse. We explain the mechanism (reachability saturates on a dense graph) and document the engineering and the methodological departures from [3].

1. Introduction

[3] show that high-scoring retrieved passages are often *distractors* — semantically similar to the query yet answer-free — and degrade the reader. We test one idea against this: filter distractors by *graph proximity* on an open knowledge graph, reranking retrieved passages by how close their entities are, on Wikidata, to the query entities. We deliberately restrict ourselves to pure topology — bare ≤ 3 -hop reachability, no relation semantics — to isolate what proximity *alone* contributes. The answer is negative: topology alone does not filter distractors. The report’s main contribution is this controlled test and the mechanistic account of *why* it fails; a by-product is a method to decide ≤ 3 -hop reachability over the entire Wikidata graph without materi-

alising hub neighbourhoods. Code {4}.

2. Related work

We adopt the corpus, retriever and reader of [3], who characterise the distractor problem in RAG. Retrieval uses Contriever over the gold-augmented Wikipedia corpus of [3]; entity linking uses the ReFinED linker of Ayoola et al. [1], one of the systems evaluated by Bast et al. [2]. Knowledge-graph approaches to RAG use the graph’s relations and semantics; we instead use only its topology. The graph is stored as HDT over a Wikidata dump.

3. Method

3.1. Baseline: dense RAG

We reproduce the dense-retrieval pipeline of [3]. The corpus is its gold-augmented Dec-2018 Wikipedia release {1}; we re-segment it on sentence boundaries and pad each passage to exactly 100 words — the raw DPR passages have ragged, subject-less tails that hurt entity linking and the reader, and since Contriever mean-pools, fixed length is only a consistency choice. Passages are Contriever-encoded and indexed with FAISS (IndexFlatIP, ~ 42 M). From NQ-open {2} we keep short (≤ 5 -token) answers — matching the extractive setup of [3] — whose question and *all* answer variants link to Wikidata entities (ReFinED); we evaluate on a curated 1000-query subset (344/1000 queries are substituted so retrieved passages carry entities; see Appendix). The top-5 passages feed Llama-2-7B (base) under a one-shot extraction prompt, and correctness is decided by an *LLM judge* (Qwen2.5-7B-Instruct, a different model to avoid self-preference) that checks whether the response contains an NQ gold answer — replacing the brittle substring exact-match of [3] (prompt and judge in the Appendix).

3.2. Contribution: KG concept-proximity reranking

We rerank the top-100 by graph proximity on Wikidata. We query a pre-built HDT dump {3} (Jan 2022, ~ 166 GB) via `pyHDT` — the public SPARQL endpoint timed out even

Email: Gianluca Como <comogianluca@gmail.com>.

Table 1. LLM-judge accuracy, all 90 rerank settings (1000 NQ-open queries). Dense baseline = 0.364. Columns: KG weight α ; rows: hop-distance k , hub-threshold t (∞ =no filter; distance 1 has no bridge, so it is threshold-invariant). Best **bold**, worst underlined.

k	t	$\alpha=.1$	.3	.5	.7	.9
1	any	.353	.344	.343	.316	.322
2	500	.362	.358	.361	.348	.354
2	1,000	.364	.361	.358	.343	.350
2	2,000	.360	.371	.357	.337	.341
2	5,000	.357	.372	.362	.337	.337
2	10,000	.364	.365	.357	.334	.329
2	∞	.354	.343	.311	.298	.295
3	500	.354	.344	.341	.329	.335
3	1,000	.346	.355	.342	.322	.310
3	2,000	.354	.349	.316	.303	.293
3	5,000	.366	.351	.310	.292	<u>.290</u>
3	10,000	.361	.362	.321	.306	.304
3	∞	.347	.333	.314	.314	.318

on a single hub count (see Appendix) — and export every direct entity-to-entity statement (Wikidata’s “truthy” wdt: predicates), 6.61×10^8 edges. For a pair (q, d) we test $\text{dist}(q, d) \leq 3$ *meet-in-the-middle*: distance 1 a set lookup, distance 2 an intersection of 1-hop neighbourhoods N_1 , distance 3 an indexed edge-probe between $N_1(q)$ and $N_1(d)$ — never materialising a hub’s neighbour list (Q30, USA: 2.35M edges). A degree threshold t drops high-degree bridges. The KG score is $s_{\text{kg}} = \text{cr} \cdot \text{pr}$, where the *connected ratio* cr is the fraction of query entities reaching ≥ 1 document entity within k hops and the *purity ratio* pr the fraction of document entities reached by any query entity. We fuse with the dense score,

$$s = (1 - \alpha) s_{\text{dense}} + \alpha s_{\text{kg}}, \quad (1)$$

with α the *KG weight* and $s_{\text{dense}} = \sigma / \max_q \sigma$ a per-query max-only rescaling of the raw Contriever score σ (the Appendix explains why not min-max). We sweep $k \in \{1, 2, 3\}$, $t \in \{500, 1k, 2k, 5k, 10k, \infty\}$, $\alpha \in \{0.1, 0.3, 0.5, 0.7, 0.9\}$: 90 rerank cells plus the dense baseline.

4. Results

Baseline is sound. Dense retrieval scores 0.364, in line with the ≈ 0.36 that [3] report for Llama-2 with Contriever-retrieved documents on NQ-open — the harness is healthy and the comparison valid.

No gain, then harm. The best cell ($\alpha=0.3$, dist 2, $t=5k$) reaches 0.372: $+0.8\text{pp} \approx 8$ queries, i.e. noise (Table 1). As α grows the KG signal dominates and accuracy falls monotonically, bottoming at 0.290 ($\alpha=0.9$, dist 3, $t=5k$).

It reshuffles, but does not help. The rerank is far from inert: at $\alpha=0.5$ it rewrites the top-5 for $>95\%$ of queries

(mean Jaccard@5 ≈ 0.12 vs. pure retrieval — barely one of five passages survives), yet accuracy stays at the baseline. It reorders aggressively without adding signal.

Threshold gradient is the mechanism. Holding $\alpha=0.9$, dist 2, accuracy decays monotonically as the hub filter loosens, from 0.354 ($t=500$) to 0.295 ($t=\infty$): admitting high-degree nodes as bridges connects everything to everything and destroys the signal.

Significance. Exact McNemar tests against the baseline, Bonferroni-corrected for the 90 comparisons (significance threshold $0.05/90 \approx 5.6 \times 10^{-4}$), flag 13 of the 90 conditions as significant — *all* worse, none better; each flips $\approx 200/1000$ answers yet nets negative.

5. Discussion and conclusions

Why it fails. On a graph as dense as Wikidata, 2–3-hop reachability *saturates*: almost every document entity is reachable from almost every query entity, so $\text{cr} \approx 1$, the score collapses onto pr alone, and discrimination is lost; the hub filter limits the damage but recovers no signal. For about 17% of queries the question entity is itself a very high-degree node (a country, a religion), where any passage is graph-close and the KG score is uninformative.

Limitations and future work. Our metric collapses relation semantics (P31 *instance-of* and P19 *place-of-birth* count alike), so it cannot screen out semantically irrelevant links. The honest next step is to use the graph’s relations and semantics — not just its bare topology.

Conclusion. Pure topological proximity does not filter distractors for dense RAG on NQ-open; the value is the explained negative result and the reusable Wikidata tooling.

Data and code

- {1} Corpus: https://huggingface.co/datasets/florin-hf/wiki_dump2018_nq_open.
- {2} Queries: https://huggingface.co/datasets/florin-hf/nq_open_gold.
- {3} Wikidata HDT dump: <https://hdt-dumps.cluster.ai.wu.ac.at/dumps/>.
- {4} Code: delivery repository (to be created).

References

- [1] Ayoola, T., Tyagi, S., Fisher, J., Christodoulopoulos, C., and Pierleoni, A. ReFinED: An efficient zero-shot-capable approach to end-to-end entity linking. In *Proceedings of the 2022 Conference of the North American*

Chapter of the Association for Computational Linguistics (NAACL): Industry Track, 2022.

- [2] Bast, H., Hertel, M., and Prange, N. A fair and in-depth evaluation of existing end-to-end entity linking systems. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2023.
- [3] Cuconasu, F., Trappolini, G., Siciliano, F., Filice, S., Campagnano, C., Maarek, Y., Tonellotto, N., and Silvestri, F. The power of noise: Redefining retrieval for RAG systems. In *Proceedings of the 47th International ACM SIGIR Conference on Research and Development in Information Retrieval (SIGIR)*, 2024.
- [4] Naveed, H., Khan, A. U., Qiu, S., Saqib, M., Anwar, S., Usman, M., Akhtar, N., Barnes, N., and Mian, A. A comprehensive overview of large language models. *arXiv preprint arXiv:2307.06435*, 2023.
- [5] Peng, B., Zhu, Y., Liu, Y., Bo, X., Shi, H., Hong, C., Zhang, Y., and Tang, S. Graph retrieval-augmented generation: A survey. *arXiv preprint arXiv:2408.08921*, 2024.

Appendix

Pipeline: notebook flow

The project runs as nine notebooks in order (the Wikidata graph itself is built separately by `scripts/pipeline/`, Layers 1–3, under WSL2):

1. `01_corpus_preparation` — download the corpus and segment it into sentence-aligned 100-word passages (~42M).
2. `02_nq_filtering` — filter NQ-open (≤ 5 -token answers + ReFiNed entities) to 31,372 queries.
3. `03_embedding` — Contriever-encode every passage and build 9 FAISS shards.
4. `04_answer_preparation` — top-100 retrieval and passage entity-linking for the 1000-query subset.
5. `05_answer_curation` — find the 344 queries whose top-100 carry no entities and pick replacements.
6. `06_apply_curation` — apply the swap, producing the `*_curated` files.
7. `07_kg_rerank` — KG reachability (cr/pr) and the Jaccard reshuffle grid.
8. `08_llm_eval` — Llama-2-7B generates answers for all 91 conditions.
9. `09_llm_judge` — Qwen2.5 judges each answer; accuracy + McNemar.

Environment

Two environments are required:

- **Windows + uv** — everything except the graph build.
- **WSL2/Ubuntu** for the HDT pipeline: `pyHDT` needs a C++ build that does not compile cleanly on Windows native; the 166 GB dump is read via `/mnt/c`.
- **faiss-gpu via conda** — no Windows pip wheel, so it is installed in miniconda and bridged into the uv venv at runtime (`add_dll_directory + sys.path.insert`; conda paths are machine-specific).
- **ReFiNed patch** — ReFiNed V1 breaks on Windows + Python 3.12 + transformers 4.48 (a Unix-only `strftime("%s")` call); `scripts/tooling/patch_refined.py` fixes it, auto-run via subprocess.
- **HuggingFace token** in `.hftoken` at the repo root, for gated models (Llama-2).

All reported numbers are reproduced from the raw outputs by `documents/verify_report_numbers.py`; per-file schemas are catalogued in `DATA.DICTIONARY.md`.

Design choices

Dense normalization: max-only, not min-max. The raw Contriever score σ (inner product, not L2-normalised) sits in a narrow per-query band; we rescale by the per-query maximum only ($s_{\text{dense}} = \sigma / \max_q \sigma$). Min-max would stretch that band to $[0, 1]$ and change the ranking *philosophy*: e.g. $A(\sigma=1.05, s_{\text{kg}}=0.2)$ vs $B(\sigma=1.00, s_{\text{kg}}=0.4)$ at $\alpha=0.5$ — max-only treats A’s 5% dense edge as negligible and picks **B**, while min-max inflates that 5% gap to 100% and flips to **A**. Max-only preserves the dense geometry and only caps it at 1.

Reader prompt. Llama-2-7B base (not chat) is a completion model: without an exemplar it continues “Answer:” with prose. We prepend a one-shot extraction example and place the question *above* the documents (vs [3], which places it after); the top-ranked passage sits last, nearest the answer slot.

Judge. The judge must differ from the generator to avoid self-preference, so we use Qwen2.5-7B-Instruct (Apache-2.0, ungated), deciding YES/NO on whether a response contains a gold answer — absorbing paraphrase and formatting that a substring match misses.

What we tried first

The final reachability method is the third design.

SPARQL → local HDT. The public Wikidata SPARQL endpoint timed out even on a minimal COUNT over a hub

(?x wdt:P31 wd:Q5); a local HDT dump queried with pyHDT gives sub-millisecond predicate-restricted counts.

Wildcard → **per-predicate export**. A wildcard triple iterator overshoot and silently dropped genuine Q–Q edges on a dump this large; exporting one predicate at a time gave exact, verifiable counts (661,471,158 edges).

BFS-N3 → **meet-in-the-middle**. Precomputing full 3-hop neighbourhoods by BFS timed out on 89% of queries (hubs explode the frontier); we instead precompute only 1-hop neighbourhoods (~93M rows) and decide ≤ 3 -hop reachability per pair by meeting in the middle.

Query curation. 344/1000 subset queries had top-100 passages with no linkable entities (KG score then undefined); we substituted them with entity-bearing queries, which biases evaluation *towards* the KG — so the negative result reads “even on KG-favourable queries it does not help”.